

ProjectMilestone_1_PuppalaSucharitha

October 9, 2022

0.1 Term Project

0.2 Project MileStone 1: Data Selection and EDA

0.3 Puppala Sucharitha

0.4 Data Science, Bellevue University

0.5 DSC550-T301 Data Mining (2231-1)

0.6 Professor : Dr. Brett Werner

0.7 Date : 10/03/2022

0.8 Introduction :

Heart disease describes a range of conditions that affect your heart. Diseases under the heart disease umbrella include blood vessel diseases, such as coronary artery disease, heart rhythm problems (arrhythmias) and heart defects you're born with (congenital heart defects), among others.

The term "heart disease" is often used interchangeably with the term "cardiovascular disease". Cardiovascular disease generally refers to conditions that involve narrowed or blocked blood vessels that can lead to a heart attack, chest pain (angina) or stroke. Other heart conditions, such as those that affect your heart's muscle, valves or rhythm, also are considered forms of heart disease.

Heart disease is one of the biggest causes of morbidity and mortality among the population of the world. Prediction of cardiovascular disease is regarded as one of the most important subjects in the section of clinical data analysis. The amount of data in the healthcare industry is huge. Data mining turns the large collection of raw healthcare data into information that can help to make informed decisions and predictions.

According to a news article, heart disease proves to be the leading cause of death for both women and men. About 610,000 people die of heart disease in the United States every year—that's 1 in every 4 deaths. Heart disease is the leading cause of death for both men and women. More than half of the deaths due to heart disease in 2009 were in men. Coronary Heart Disease (CHD) is the most common type of heart disease, killing over 370,000 people annually.

This makes heart disease a major concern to be dealt with. But it is difficult to identify heart disease because of several contributory risk factors such as diabetes, high blood pressure, high cholesterol, abnormal pulse rate, and many other factors. Due to such constraints, scientists have turned towards modern approaches like Data Mining and Machine Learning for predicting the disease.

Machine learning (ML) proves to be effective in assisting in making decisions and predictions from the large quantity of data produced by the healthcare industry.

This project focuses on to classify / predict whether a patient is prone to heart disease depending on multiple attributes.

In this project the following steps are involved :

- Overview of the Dataset.
- Exploratory Data Analysis.
- Data Preparation.
- Identifying the features that are most influential for having Heart Disease.
- Build different models to predict Heart Disease.

This dataset contains 11 features that can be used to predict a possible heart disease.

1. Age : age of the patient [years]
2. Sex : sex of the patient [M: Male, F: Female]
3. ChestPainType : chest pain type [TA: Typical Angina, ATA: Atypical Angina, NAP: Non-Anginal Pain, ASY: Asymptomatic]
4. RestingBP : resting blood pressure [mm Hg]
5. Cholesterol : serum cholesterol [mm/dl]
6. FastingBS : fasting blood sugar [1: if FastingBS > 120 mg/dl, 0: otherwise]
7. RestingECG : resting electrocardiogram results [Normal: Normal, ST: having ST-T wave abnormality (T wave inversions and/or ST elevation or depression of > 0.05 mV), LVH: showing probable or definite left ventricular hypertrophy by Estes' criteria]
8. MaxHR : maximum heart rate achieved [Numeric value between 60 and 202]
9. ExerciseAngina : exercise-induced angina [Y: Yes, N: No]
10. Oldpeak : oldpeak = ST [Numeric value measured in depression]
11. ST_Slope : the slope of the peak exercise ST segment [Up: upsloping, Flat: flat, Down: downsloping]
12. HeartDisease : output class [1: heart disease, 0: Normal]

The Heart Failure Prediction dataset is collected from Kaggle.com and the link to the dataset is given below.

<https://www.kaggle.com/datasets/fedesoriano/heart-failure-prediction>

References :

- Shubhankar Rawat. "HeartDisease Prediction." Medium, Towards Data Science, 10 Aug 2019, <https://medium.com/towards-data-science/heart-disease-prediction-73468d630cfc>
- fedesoriano. (September 2021). Heart Failure Prediction Dataset. Retrieved [Date Retrieved] from <https://www.kaggle.com/fedesoriano/heart-failure-prediction>.

0.9 Milestone -1 : Data Selection and Exploratory Data Analysis(EDA)

In the Milestone-1, I would like to get an overview of the data set collected and the Exploratory Data Analysis of the the dataset.

In this Milestone-1, I would like to know the following points:

- Distribution of the different features in the data set on the HeartDisease.
- To know the counts of the categorical features available in the dataset.
- Are there any duplicated, missing values and outliers in the dataset.

```
[1]: # Importing all the necessary libraries.
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib as mpl
import matplotlib.pyplot as plt
from scipy.stats import skew
from sklearn import linear_model
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import Lasso
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import mean_absolute_error
from sklearn.metrics import mean_squared_error
from sklearn.metrics import r2_score
%matplotlib inline
import warnings
warnings.filterwarnings('ignore')
```

```
[2]: # Loading the heart dataset which is named heart.csv
heartdf = pd.read_csv('heart.csv')
```

```
[3]: # Getting the first 5 rows of the heart dataset.
heartdf.head()
```

```
[3]:
```

	Age	Sex	ChestPainType	RestingBP	Cholesterol	FastingBS	RestingECG	MaxHR	\
0	40	M	ATA	140	289	0	Normal	172	
1	49	F	NAP	160	180	0	Normal	156	
2	37	M	ATA	130	283	0	ST	98	
3	48	F	ASY	138	214	0	Normal	108	
4	54	M	NAP	150	195	0	Normal	122	

	ExerciseAngina	Oldpeak	ST_Slope	HeartDisease
0	N	0.0	Up	0
1	N	1.0	Flat	1
2	N	0.0	Up	0
3	Y	1.5	Flat	1
4	N	0.0	Up	0

```
[4]: # Getting the shape of the heart dataset
heartdf.shape
```

```
[4]: (918, 12)
```

```
[5]: # Getting the size of the dataset
heartdf.size
```

```
[5]: 11016
```

```
[6]: # Information about the dataset variables.
heartdf.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 918 entries, 0 to 917
Data columns (total 12 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   Age                   918 non-null    int64
 1   Sex                   918 non-null    object
 2   ChestPainType         918 non-null    object
 3   RestingBP             918 non-null    int64
 4   Cholesterol           918 non-null    int64
 5   FastingBS             918 non-null    int64
 6   RestingECG           918 non-null    object
 7   MaxHR                 918 non-null    int64
 8   ExerciseAngina        918 non-null    object
 9   Oldpeak               918 non-null    float64
10   ST_Slope              918 non-null    object
11   HeartDisease          918 non-null    int64
dtypes: float64(1), int64(6), object(5)
memory usage: 86.2+ KB
```

```
[7]: # Checking for null values
heartdf.isna().sum()
```

```
[7]: Age                0
     Sex                0
     ChestPainType     0
     RestingBP         0
     Cholesterol       0
     FastingBS         0
     RestingECG        0
     MaxHR             0
     ExerciseAngina    0
     Oldpeak           0
     ST_Slope          0
     HeartDisease      0
dtype: int64
```

```
[8]: # Checking if any duplicated values present in the data set.
heartdf.duplicated()
```

```
[8]: 0      False
      1      False
      2      False
      3      False
      4      False
      ...
      913    False
      914    False
      915    False
      916    False
      917    False
      Length: 918, dtype: bool
```

```
[9]: # Getting the number of unique columns in the data set
      heartdf.nunique()
```

```
[9]: Age                50
      Sex                2
      ChestPainType      4
      RestingBP          67
      Cholesterol        222
      FastingBS          2
      RestingECG         3
      MaxHR             119
      ExerciseAngina     2
      Oldpeak           53
      ST_Slope           3
      HeartDisease       2
      dtype: int64
```

Observations of the dataset:

- The dataset has 918 rows and 12 columns.
- We can see that the dataset contains both numerical and categorical data.
- In the initial review of the dataset we can see there are no null values.
- Extracted the number of unique values for each column.
- Number of unique values for each column is extracted above and will be reviewed during the EDA section.

0.9.1 Exploratory Data Analysis.

```
[10]: # Initially let's use the describe() to get an idea on the dataset.
      heartdf.describe()
```

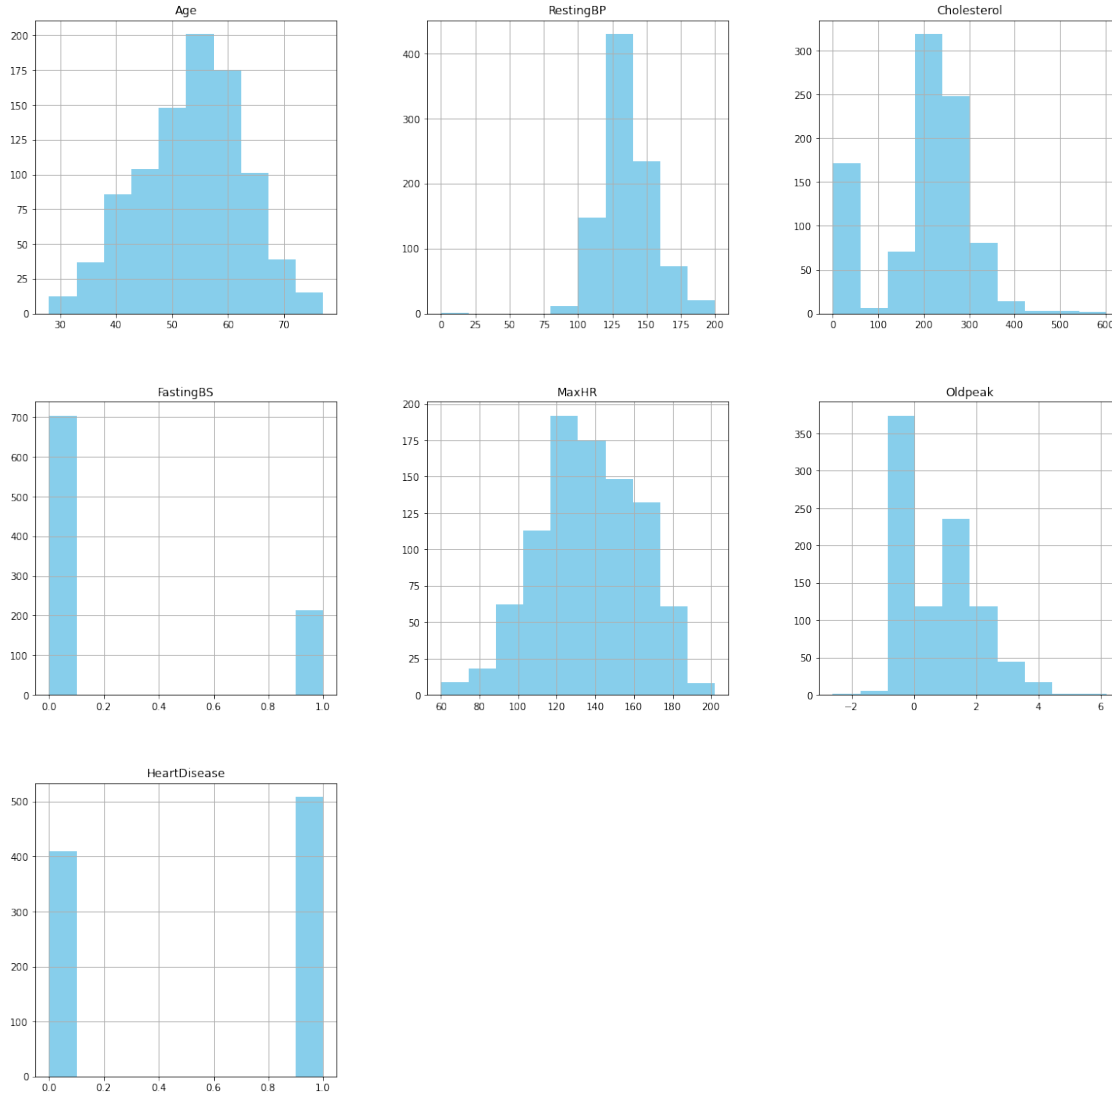
```
[10]:
```

	Age	RestingBP	Cholesterol	FastingBS	MaxHR	\
count	918.000000	918.000000	918.000000	918.000000	918.000000	
mean	53.510893	132.396514	198.799564	0.233115	136.809368	
std	9.432617	18.514154	109.384145	0.423046	25.460334	

min	28.000000	0.000000	0.000000	0.000000	60.000000
25%	47.000000	120.000000	173.250000	0.000000	120.000000
50%	54.000000	130.000000	223.000000	0.000000	138.000000
75%	60.000000	140.000000	267.000000	0.000000	156.000000
max	77.000000	200.000000	603.000000	1.000000	202.000000

	Oldpeak	HeartDisease
count	918.000000	918.000000
mean	0.887364	0.553377
std	1.066570	0.497414
min	-2.600000	0.000000
25%	0.000000	0.000000
50%	0.600000	1.000000
75%	1.500000	1.000000
max	6.200000	1.000000

```
[11]: # plot histograms for each numerical variable
heartdf.hist(figsize = (20, 20),color='skyblue')
plt.show()
```



Observations from the histogram plots of the dataset.

- From the above plots we can see some features are skewed.


0.9.2 Visualization of the target variable

```
[12]: # Let's plot a pie plot and count plot of the target variable 'HeartDisease'
l = list(heartdf['HeartDisease'].value_counts())
circle = [l[1] / sum(l) * 100, l[0] / sum(l) * 100]
# Setting the colors of the plots.
colors = ['green', 'orange']

# Plotting the pie plot.
fig, ax = plt.subplots(nrows = 1, ncols = 2, figsize = (20, 5))
```

```
plt.subplot(1,2,1)
plt.pie(circle,labels = ['No Heart Disease','Heart Disease'],autopct='%1.
    ↳1f%%',startangle = 90,explode = (0.1,0),colors = colors)
plt.title('Distribution of the target variable - Heart Disease in %'); # Title of the plot.

# Plotting the count plot.
plt.subplot(1,2,2)
sns.countplot('HeartDisease',data = heartdf,palette = colors)
plt.title('Distribution of the target variable - Heart Disease');# Title of the plot
plt.show()
```




Observations from the target variable pie plot and count plot.

- From the above visualizations we can see that the dataset is evenly balanced.
- We can see from the pie plot that 55.3% are having Heart Disease and 44.7% are not having any Heart Disease.

0.9.3 Visualization of the target variable with sex.

```
[13]: #Let's check for the distribution of HeartDisease by sex.
sns.countplot(data=heartdf,x='Sex' , palette=['skyblue','pink']);
#Title of the plot
plt.title('HeartDisease Distribution by Sex')
# X- label
plt.xlabel('Sex')
# Y- label
plt.ylabel('Count')
```


```
[13]: Text(0, 0.5, 'Count')
```

Observations of the target variable by sex.

- Here we can see that males are highly prone to have HeartDisease.

```
[14]: # Let's plot a pie plot for the distribution of heart disease by sex.
plt.rcParams.update({'font.size': 10})
ax=heartdf['Sex'].value_counts().plot.pie(explode=[0.5,0.5],autopct='%1.
    ↪2f%%',shadow=True );
# Plot title.
ax.set_title(label = "Heartdisease Distribution by Sex");
plt.axis('off');
```



Observations of the pie plot :

- From the above pie plot we see that about 78.98% of males have heart disease and 21.02% of females have HeartDisease.
- We can say that Males are approximately 3 times more likely to have HeartDisease than females.

0.9.4 Dividing the Categorical and Numerical variables from the dataset.

```
[15]: # Let's divide the categorical and numerical variables from the dataset.
col = list(heartdf.columns)
categ_features = []
numer_features = []
for i in col:
    if heartdf[i].dtype == np.object:
        categ_features.append(i)
    else:
        numer_features.append(i)


print('Categorical Features of the heart dataset :',*categ_features)
print('Numerical Features of the heart dataset :',*numer_features)
```



Categorical Features of the heart dataset : Sex ChestPainType RestingECG
ExerciseAngina ST_Slope



Numerical Features of the heart dataset : Age RestingBP Cholesterol FastingBS
MaxHR Oldpeak HeartDisease



0.9.5 Plotting Numerical features with Target Variable i.e HeartDisease.

```
[16]: # Plotting numerical features with target variable.  
# With these plots we compare between our target and our numerical features.  
#colors = ['green','orange']  
for i in numer_features:  
    plt.figure(figsize=(10,5))  
    sns.countplot(x=i, data=heartdf, hue='HeartDisease')  
    plt.legend(['Normal', 'Heart Disease'])  
    plt.title(i + " vs Heart Disease")# Plot title  
    plt.show()
```










Observations:

- From the above count plot of Age vs HeartDisease most of the heart disease Patients have age between 55 and 65
- From the above count plot of Fasting Blood Sugar vs HeartDisease , we see that the persons having FastingBS and positive HeartDisease counts upto 150.


- From the above graphs we can see that the distribution of the features Rest-ingBP,Cholesterol,MaxHR and Oldpeak are not clearly understandable because of the presence of many unique data points in the features.
- For more clear understanding let us scale the individual values of these features which brings the data points to a constant value representing a range of values.
- Here I have divided the data points of the numerical features by 5 or 10 and assigned the quotient value as the representative of the constant data point.

```
[17]: # Scaling the 'RestingBP'
colors = ['green','orange']
heartdf['RestingBP_Group'] = [ int(i / 5) for i in heartdf['RestingBP']]
restingBP_group = heartdf['RestingBP_Group']
sns.countplot(restingBP_group,data = heartdf,hue = "HeartDisease",palette = colors)
title = 'RestingBP_Group vs Heart Disease' # Plot title
plt.title(title);
```




```
[18]: # Scaling of 'Cholesterol'
heartdf['Cholesterol_Group'] = [ int(i / 10) for i in heartdf['Cholesterol']]
cholesterol_group = heartdf['Cholesterol_Group']
plt.figure(figsize=(10,5))
```


```
sns.countplot(cholesterol_group,data = heartdf,hue = "HeartDisease",palette = colors)
title = 'Cholesterol_Group vs Heart Disease'# Plot Title.
plt.title(title);
```



```
[19]: # Scaling of 'MaxHR'
heartdf['MaxHR_Group'] = [ int(i / 5) for i in heartdf['MaxHR']]
maxhr_group = heartdf['MaxHR_Group']
sns.countplot(maxhr_group,data = heartdf,hue = "HeartDisease",palette = colors)
title = 'MaxHR_Group vs Heart Disease' # Plot Title
plt.title(title);
```

```
[20]: # Scaling of 'Oldpeak'
heartdf['Oldpeak_Group'] = [ int( (i*10) / 5) for i in heartdf['Oldpeak']]
oldpeak_group = heartdf['Oldpeak_Group']
sns.countplot(oldpeak_group,data = heartdf,hue = "HeartDisease",palette =_
    ↪colors)
title = 'Oldpeak_Group vs Heart Disease'# Plot title.
plt.title(title);
```



Observations of count plots of :






- From the graph of scaled RestingBP, we can say that the values 95 to 170 ((195)- (345))are mostly prone to have HeartDisease.
- From the graph of scaled Cholesterol, we can say that the values 160 to 340 ((1610)- (3410))are prone to have HeartDisease.
- From the graph of scaled MaxHR, we can say that the values 70 to 180 ((145) - (365)) are mostly prone to have HeartDisease.
- From the graph of sccaled Oldpeak, we can say that the values 0 to 8 i.e 0 to 4((05/10) - (85/10)) are mostly prone to have HeartDisease.

```
[21]: # Drop the grouped columns that are created.
heartdf = heartdf.drop(columns = [
    ↪ ["Oldpeak_Group", "RestingBP_Group", "Cholesterol_Group", "MaxHR_Group"]])
```

0.9.6 Plotting Categorical Features with the Target Variable i.e. Heart Disease.

```
[22]: # Plotting categorical features with the target variable.
# here we compare between target and Categorical features
colors = ['green', 'orange']
fig, ax = plt.subplots(nrows = 3, ncols = 1, figsize = (10,15))
for i in range(len(categ_features)):
```

```
plt.subplot(3,2,i+1)
sns.countplot(categ_features[i],data = heartdf,hue = "HeartDisease",palette_
↪= colors)
title = categ_features[i] + ' vs HeartDisease'
plt.title(title);
```



Observations of Count plots: The following observations are made from the above plots.

- From the above 'Sex vs HeartDisease' plot we can say that the Male population has more prone to HeartDisease and Female population are less prone to HeartDisease.
- The 'ChestPainType vs HeartDisease' plot has four types they are - TA: Typical Angina, ATA: Atypical Angina, NAP: Non-Anginal Pain, ASY: Asymptomatic among these four ASY type has more positive HeartDisease cases.
- From the 'RestingECG vs HeartDisease' plot we see that all the three types are having Heart Disease.
- From the 'ExerciseAngina vs HeartDisease' plot we see that the persons having ExerciseAngina are more prone to have HeartDisease.
- From the 'ST_Slope vs HeartDisease' plot we can see that the persons having flat slope are more prone to have HeartDisease.

0.9.7 Visualizations of Categorical Features vs Positive Heart Disease Cases :

```
[23]: # Getting the postive heart disease data for the categorial features.

# Sex feature
sex = heartdf[heartdf['HeartDisease'] == 1]['Sex'].value_counts()
sex = [sex[0] / sum(sex) * 100, sex[1] / sum(sex) * 100]

# ChestPainType
chestpain = heartdf[heartdf['HeartDisease'] == 1]['ChestPainType'].
    ↳value_counts()
chestpain = [chestpain[0] / sum(chestpain) * 100,chestpain[1] / sum(chestpain)↳
    ↳* 100,chestpain[2] / sum(chestpain) * 100,chestpain[3] / sum(chestpain) *↳
    ↳100]

# RestingECG
restingecg = heartdf[heartdf['HeartDisease'] == 1]['RestingECG'].value_counts()
restingecg = [restingecg[0] / sum(restingecg) * 100,restingecg[1] /↳
    ↳sum(restingecg) * 100,restingecg[2] / sum(restingecg) * 100]

# ExerciseAngina
exerangina = heartdf[heartdf['HeartDisease'] == 1]['ExerciseAngina'].
    ↳value_counts()
exerangina = [exerangina[0] / sum(exerangina) * 100,exerangina[1] /↳
    ↳sum(exerangina) * 100]

# ST_Slope
stslope = heartdf[heartdf['HeartDisease'] == 1]['ST_Slope'].value_counts()
stslope = [stslope[0] / sum(stslope) * 100,stslope[1] / sum(stslope) *↳
    ↳100,stslope[2] / sum(stslope) * 100]
```

```
[24]: # Plotting Pie plots for the Categorical Variables with positive HeartDisease.

ax,fig = plt.subplots(nrows = 3,ncols = 1,figsize = (15,15))
colors = ['skyblue','pink']

# Pie plot of HeartDisease by Sex.
plt.subplot(3,2,1)
plt.pie(sex,labels = ['Male','Female'],autopct='%1.1f%%',startangle = 90,explode = (0.1,0),colors = colors)
plt.title('Distribution of Positive HeartDisease by Sex');# Plot title


# Pie plot of HeartDisease by ChestPainType.
plt.subplot(3,2,2)
plt.pie(chestpain,labels = ['ASY', 'NAP', 'ATA', 'TA'],autopct='%1.1f%%',startangle = 90,explode = (0,0.1,0.1,0.1))
plt.title('Distribution of Positive HeartDisease by ChestPainType'); # Plot title

# Pie plot of HeartDisease by RestingECG.
plt.subplot(3,2,3)
plt.pie(restingecg,labels = ['Normal','ST','LVH'],autopct='%1.1f%%',startangle = 90,explode = (0,0.1,0.1))
plt.title('Distribution of Positive HeartDisease by RestingECG'); #Plot Title


#Pie plot of HeartDisease by ExerciseAngina.
plt.subplot(3,2,4)
plt.pie(exerangina,labels = ['Angina','No Angina'],autopct='%1.1f%%',startangle = 90,explode = (0.1,0))
plt.title('Distribution of Positive HeartDisease of ExerciseAngina'); # Plot Title

#Pie plot of HeartDisease by ST_Slope.
plt.subplot(3,2,5)
plt.pie(stslope,labels = ['Flat','Up','Down'],autopct='%1.1f%%',startangle = 90,explode = (0,0.1,0.1))
plt.title('Distribution of Positive HeartDisease by ST_Slope'); # Plot Tile
```


Distribution of Positive HeartDisease by Sex




Distribution of Positive HeartDisease by ChestPainType




Distribution of Positive HeartDisease by RestingECG



Distribution of Positive HeartDisease of ExerciseAngina



Distribution of Positive HeartDisease by ST_Slope



Observations of Pie plots : The above plots are plotted for more clear idea about the percentage of impact of the categorical feature on positive HeartDisease

- From the pie plot of 'Sex Vs Positive HeartDisease' we see that about 90% of the patients are male.
- From the pie plot of 'ChestPainType vs Positive HeartDisease' we see that about 77.2% of the patient with HeartDisease are having ASY type of Chest Pain.
- From the pie plot 'RestingECG vs Positive HeartDisease' we see that the patients having HeartDisease have RestingECG normal level which is 56.1%
- From the pie plot 'ExerciseAngina' we see that the patients having HeartDisease , have the

problem of Exercise Induced Angina of about 62.2%

- From the pie plot 'ST_Slope vs Postivie HeartDisease' we see that about 75% of patients having HeartDisease have a flat ST_Slope.

[25]: *# Let's use pairplot() function from seaborn to understand the relationship
↪ between all features.*

```
sns.pairplot(heartdf, hue="HeartDisease");
```



[26]: *# Let's plot a correlation heatmap of the heart disease dataset.*

```
# Calculate the correlation coefficient with corr().
```

```
corr_number = heartdf.corr()
```


```
# Create the heatmap for the correlation coefficients calculated above.
```



```
fig, ax = plt.subplots(1, 1, figsize=(10,7), tight_layout = True)
sns.heatmap(corr_number, annot = True, cmap = 'flare')

# Title of the plot
plt.title('Correlation Heatmap of Heart Disease dataset')
```

[26]: Text(0.5, 1.0, 'Correlation Heatmap of Heart Disease dataset')




Observations of Pairplot and Correlation Heat Map:

- From the pair plot above we can see that some features are correlated with the target variable. In individuals who are older and having high RestingBP are more prone to HeartDisease.
- From the above correlation heat map we see some features are positively correlated with HeartDisease and some features are negatively correlated, but the correlation seems to be very low.
- It is better to create dummy variables for the categorical variables of the dataset and regenerate the correlation matrix of the dataset.
- As per the above correlation heat map we can say that Age, RestingBP are having good correlation with the HeartDisease.

0.9.8 Box plot to identify Outliers :

```
[27]: # let's plot a boxplot for the heart dataset.  
plt.figure(figsize=(20,15))  
sns.boxplot(data = heartdf)  
# Plot title  
plt.title('Boxplot of the Heart dataset')
```

```
[27]: Text(0.5, 1.0, 'Boxplot of the Heart dataset')
```



Observations of boxplot :

- From the above boxplot we see that there are outliers present in the features Age, RestingBP, Cholesterol, Oldpeak and MaxHR.

0.9.9 Conclusions from Milestone -1 : Data Selection and EDA.

- The dataset has 918 rows and 12 columns.
- The target variable “HeartDisease” is well distributed.
- There are no missing values and duplicated data in the dataset.

- The observations from the plots of patients having positive HeartDisease with different variables in the dataset says that the males are more in number, ChestPainType ASY is more dominant, FastingBS is less, RestingECG is low, ST_Slope is flat.
- The people of age above 50 years, having RestingBP between 95 - 170 , having Cholesterol ranging between 160 - 340 are more prone to HeartDisease.
- We can see some features having a good correlation with the HeartDisease like Age, Sex, but the correlation strength is very low.
- In the data preparation step it is recommended to create dummy variables for the categorical variables and regenerate the correlation matrix for collinearity check.
- Ethical CONsiderations - The ethical concerns that are to be considered in handling the data are privacy, confidentiality, honesty and fairness. While handling the data I have taken care that the data is not biased towards any factors causing heartdisease. All the data is validated to ensure that the facts are not misrepresented in the visualizations.

[]: